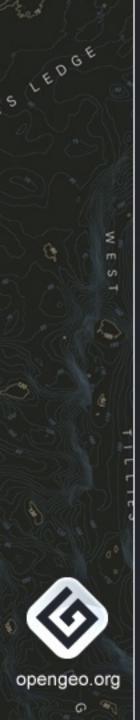


Comparing the Performance of Open Source Web Map Servers

Andrea Aime - OpenGeo

Justin Deoliveira - OpenGeo



Outline

- Test environment and setup
- WMS tests
 - vector data (plain and thematic maps)
 - raster data
- Tile cache tests
- WFS tests

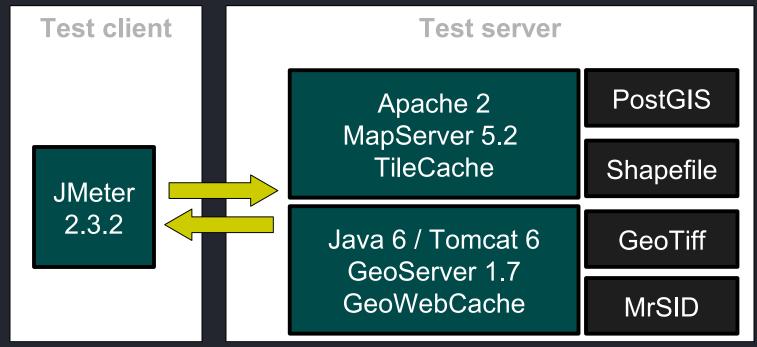


Test objectives

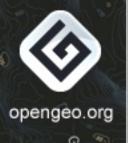
- Find out the best performing configurations (setup matters!) and making sure the servers are doing the same job
- Assess performance levels on various common tasks not tested at the FOSS4G 2007 performance comparison:
 - thematic maps
 - raster data
 - tile caching
 - WFS
- Provide feedback to the respective developer and user communities

Tes

Test enviroment



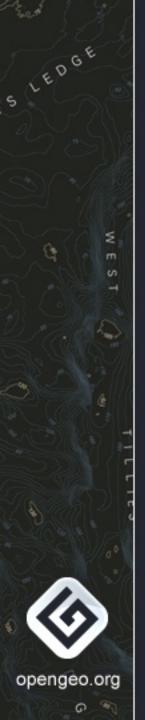
Dual core 2.1Ghz, 3GB RAM, Ubuntu 8.10



100Mbit network

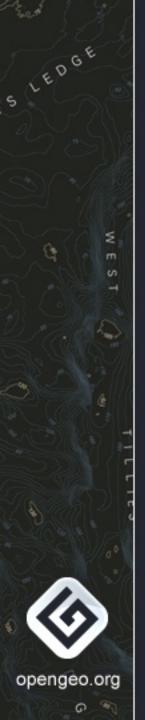






Test methodology

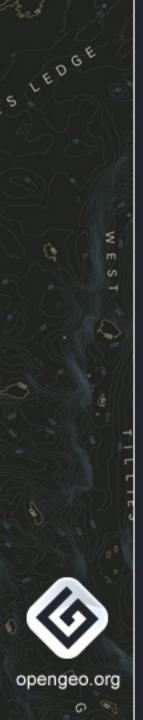
- Scale up: 1, 10, 20, 40 client threads
- Tests are repeated multiple times (cold testing does not make sense with Java servers, JIT needs time)
- Make sure the network is not playing the bottleneck (100MBit being a bottleneck? yeah! The theoretical max speed is "only" 12.5MByte/s you know?)
- Compare result quality, not just times
 - this actually uncovered bugs/potential for improvement in both servers



Setting up MapServer 5.2

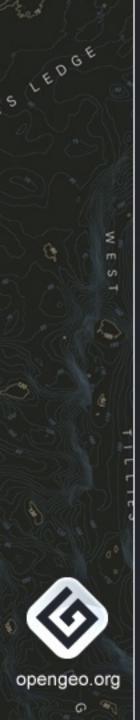
- Rebuild from sources:
 - to enable fastcgi*
 - to get MrSID support in GDAL and use the same GDAL version as GeoServer (GDAL 1.4.2)
- Move used EPSG codes at the top of the Proj / usr/share/proj/epsg file (this is crucial to get best performance)
- Use server based fastcgi, 20 mapserver processes always ready to serve requests, each one holding a single database connection (dynamic fastcgi did not work)

^{*:} very important to get good performance against PostGIS, a 3x speedup factor was measured



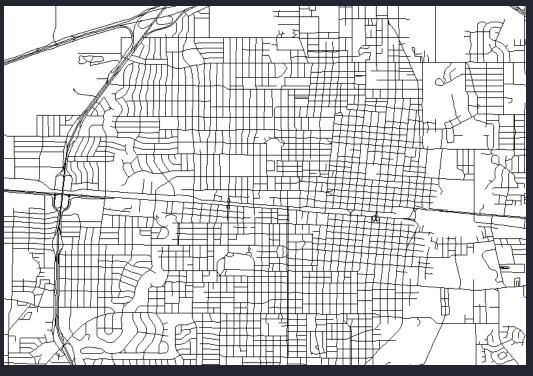
Setting up GeoServer 1.7

- Install Sun Java 6, JAI, JAI Image I/O, and ImageIO-Ext native extensions
- Install latest Tomcat release
- Disable gzip compression filters (useful on a real internet connection, but not in a local network case)
- Make sure exactly 20 threads answers requests, and that exactly 20 dbms connections are used, to match MapServer setup
- No Java VM tweaking needed, the machine is recognized as server class, the VM self tunes in this case (Linux is marked as a server OS, dual core, has more than 2GB RAM)



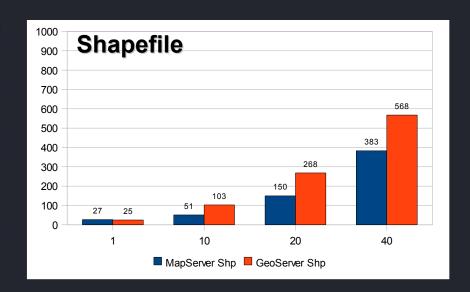
Test #1: PostgGIS vs Shapefiles

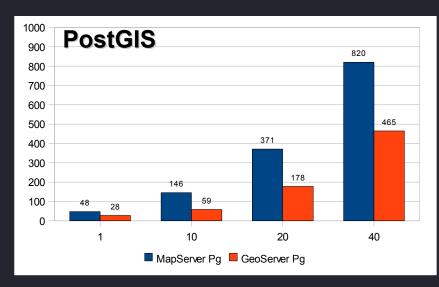
- 3 million roads in Texas
- Each test paints 1000 roads using a small bbox (512x512 tiles), minimal styling, no antialiasing



data preview

#1: Performance results





- Response time
 vs concurrent
 requests growing
 from 1 to 40
- A similar test was run one year go, compared to it:
 - MapServer improved shapefile a lot
 - GeoServer improved both



#1: Test output samples



No significant differences in output quality

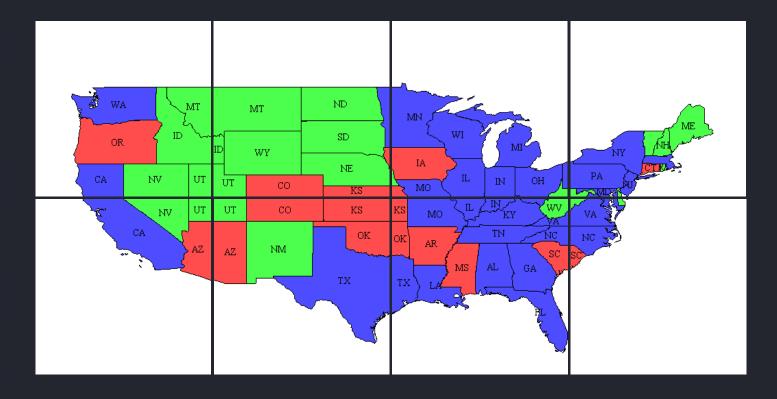


LEDGE

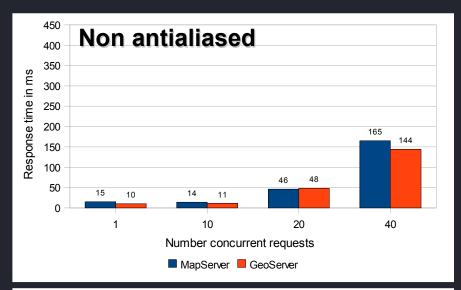
LEDGE opengeo.org

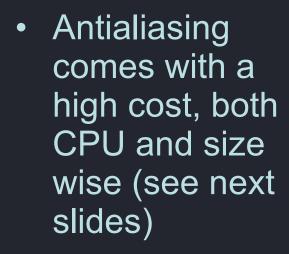
#2: Thematic map

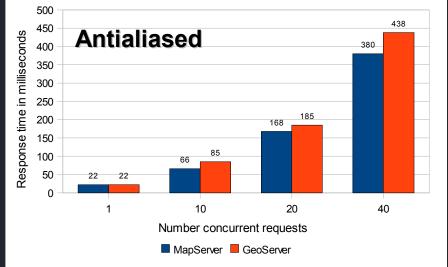
- A USA thematic map, classified by population, with labels
- 8 256x256 tiles, mimickying OpenLayers requests



#2: Performance results





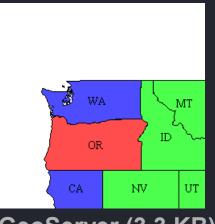


 Might be a good idea to turn it off in resource constrained environments (network/cpu)

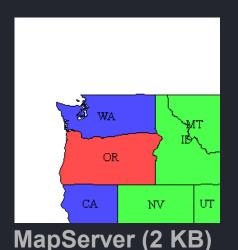


opengeo.org

#2: Test outputs (non aa)



GeoServer (3.3 KB)



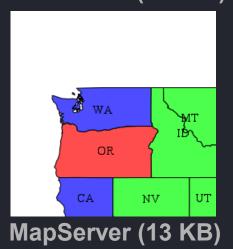
- Linework and filling area basically identical
- MapServer could use some more attention labelling (being worked on, the fix will be released in MapServer 5.4)
- Both images are very small, MapServer shows better compression rate

opengeo.org

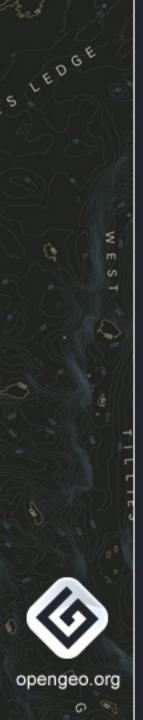
#2: Test output (antialiased)



GeoServer (9.6 KB)



- GeoServer lines too thin for 1 pixel width (being worked on, fix will be released in version 1.7.1)
- GeoServer shows better compression ratio on antialiased data
- No matter what server is being considered, the output size is 4-6 times larger than non antialiased case

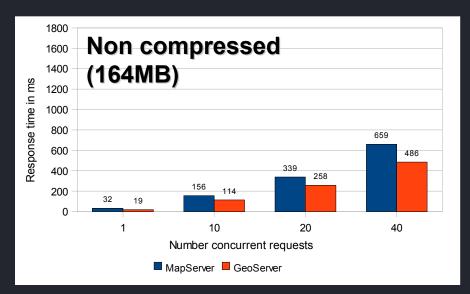


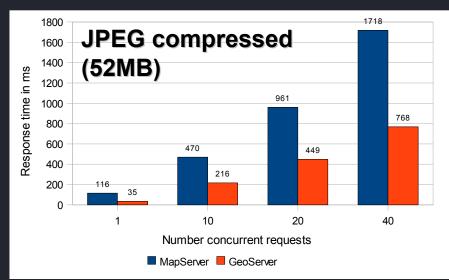
#3: Single GeoTIFF

- A single GeoTIFF file, 6800x6000:
 - Tiled, embedded overviews
- Extracting 256x256 tiles, using 1300+ different bboxes to make sure there is no caching



#3: performance results

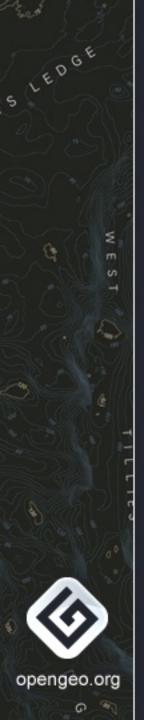




- JPEG compressing tiles reduces space consumption by 3 times
- But, it slows down decoding visibly (much more so for MapServer)



LEDGE

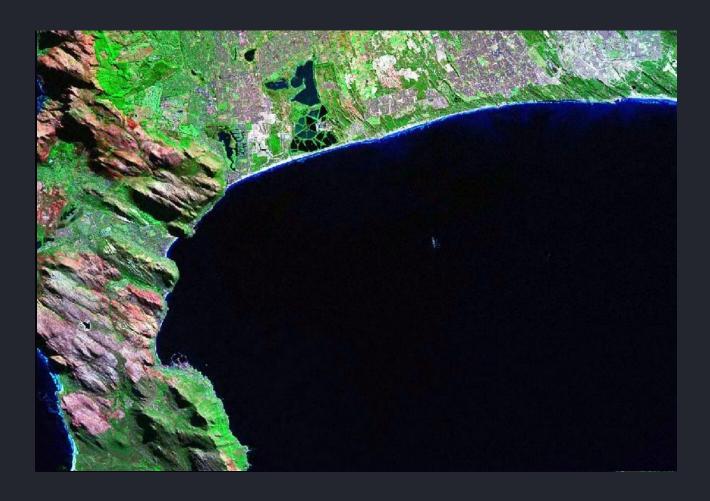


#4: MrSID data

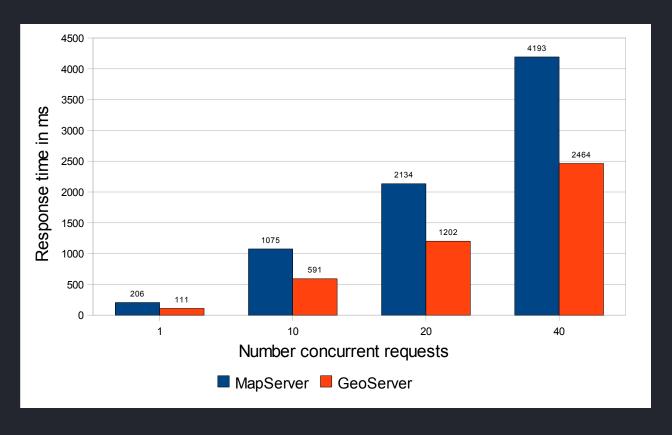
- Landsat, 47637x39167, 206MB, covering the southern part of South Africa
- https://zulu.ssc.nasa.gov/mrsid/
- Both servers using MrSID SDK 3.3 and GDAL
- Extracting 256x256 tiles, jpeg compressed, using 1300+ different bboxes to make sure there is no caching

opengeo.org

#4: a look at the data

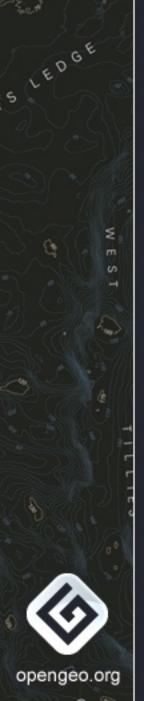


#4: performance results



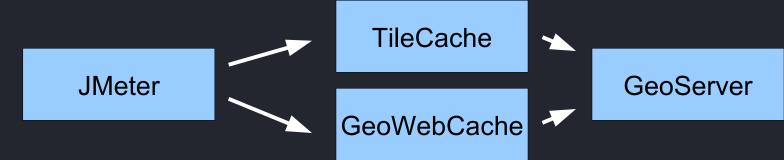
 HP: MapServer is opening the MrSID file for each request, GeoServer is keeping it open (opening and reading MrSID metadata is expensive)





#5: Tile caches

- Testing
 - TileCache 2.04 + mod_python
 - GeoWebCache 1.0-beta0
- Same backend, GeoServer 1.7.0, WMS requests
- Dataset: the MrSid file of the previous test
- Tiles reprojected to WGS84 (jpeg encoding)
- Tile caches are pre-seeded

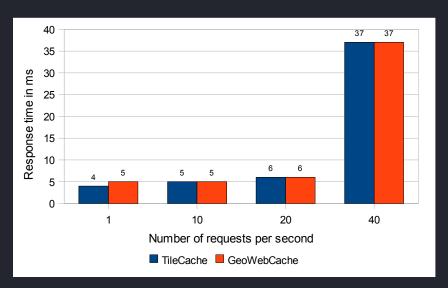


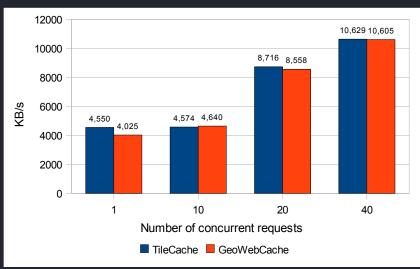
opengeo.org

#5: a look at the data



#5: performance results

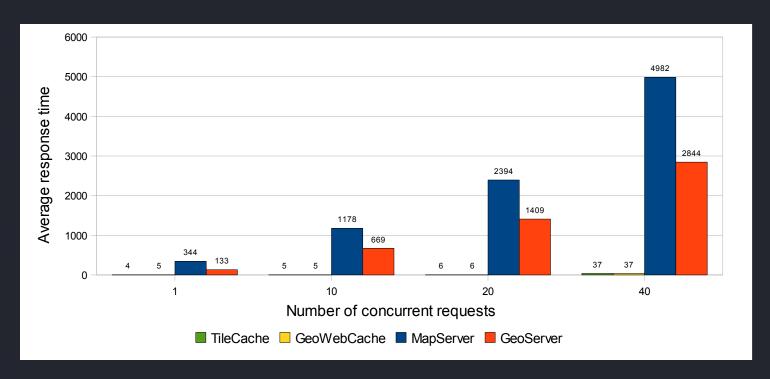




- Virtually identical results, 450+ tiles per second
- The 100MBit connection is saturated, running the test directly on the server delivers a throughput of over 17MB/s

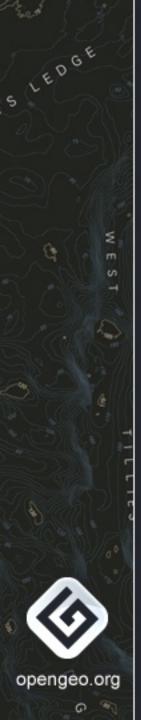


#5: quantifying the tile caching gain



- Made the same requests against MapServer and GeoServer un-cached
- In this case, the benefit is around 100 times
- The gain won't be as much if you consider simple vector rendering



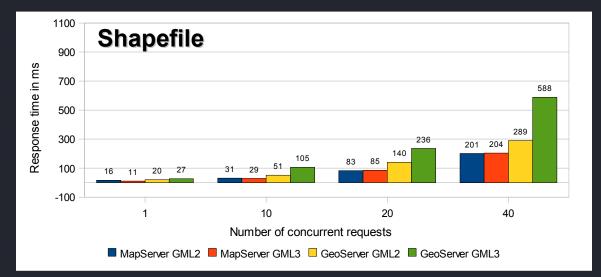


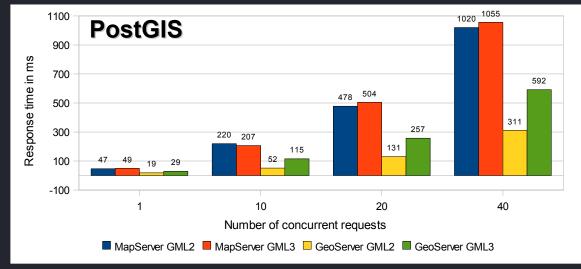
#6: WFS requests

- WFS requests with BBOX filtering
- Using the same tiles and data as the thematic map benchmark
- Shapefile and PostGIS backends
- WFS 1.0 and WFS 1.1 GetFeature

```
<wfs:FeatureCollection xsi:schemaLocation="http://www.openplans.org/topp.http://localhost:8080/geoserver/wfs?service=WFS&</p>
version=1.0.0&request=DescribeFeatureType&typeName=topp:states.http://www.opengis.net/wfs.http://localhost.8080/geoserver/schemas
/wfs/1.0.0/WFS-basic.xsd*>
- <gml: boundedBy>
  - <gral: Box srsName="http://www.opengis.net/gml/srs/epsg.xml#4326">
       <gml:coordinates decimal="."cs=","ts="">-96.640396,24.955967 -71.870476,48.173923/gml:coordinates>
    </gml:Box>
  </gral:boundedBy>
  <gml:featureMember>
  - <topp:states fid="states.1">
     - <topp:the_geom>
       - <gml: MultiPolygon srsName="http://www.opengis.net/gml/srs/epsg.xml#4326">
          - <gml:polygonMember>
            - <gml:Polygon>
              - <gral: outerBoundaryIs>
                 - <gml:LinearRing>
                   - <gml:coordinates decimal="." cs="," ts=" ">
                       -88.071564,37.51099 -88.087883,37.476273 -88.311707,37.442852 -88.359177,37.409309
                       -88.419853,37.420292 -88.467644,37.400757 -88.511322,37.296852 -88.501427,37.257782
                       -88.450699,37.205669 -88.422516,37.15691 -88.45047,37.098671 -88.476799,37.072144
                       -89.284233,37.091244 -89.303291,37.085384 -89.3097,37.060909 -89.264244,37.027733
```

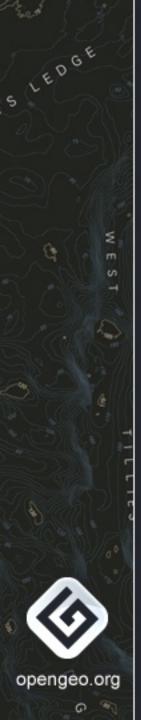
#6: performance results





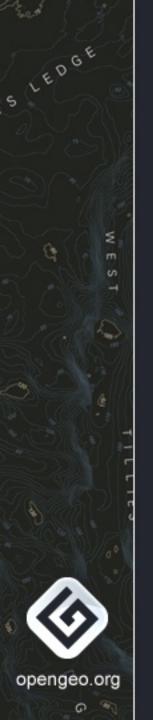
- MapServer is very fast with shapefiles, not as much with PostGIS
- GeoServer
 GML3
 output
 could use
 some work





Summary

- Both servers are very fast (most of the test deliver more than 50 responses per second... on a notebook!)
- But proper configuration is crucial to get good results
- As we have seen, both server can still work to get better, as the relative differences in results do show





QUESTIONS?